

# A Shallow Introduction to Deep Learning Tutorial

Ghent University, LT3 Group  
Thu 30 June 2016

[mike.kestemont@uantwerp.be](mailto:mike.kestemont@uantwerp.be)  
University of Antwerp, Department of Literature  
[www.mike-kestemont.org](http://www.mike-kestemont.org)

# Aim

- Intuitive introduction to Deep Learning
  - Almost no mathematics or calculus
- Lots of hands-on coding in notebooks
- Show how (surprisingly) simple it can be
- Theory in morning; Practice in afternoon

# Requirements

- Python 2.7+ (All code compatible with Python 3); Anaconda?
- External (apart from Jupiter for notebooks)
  - Scikit-learn
  - Numpy
  - Matplotlib
- DL specific (bleeding edge versions from Github)
  - TensorFlow
  - Theano
  - Keras


# Disclaimer

- Training as philologist
- Hobby that got out of hand...
- Not good at math and calculus
- Most of what I know through self-study



# Model

- System that takes an input to produce a certain output
- Has set of parameters  $\Theta$  that can be adjusted to produce a different output
- Model = System = Function  $f$ 
  - $f_{\Theta}(\text{input}) \rightarrow \text{output}$



Magic Box

# ‘Neural’ networks

- Historically **inspired** by working human brain
  - Exaggerated in media...
  - But interesting parallels
- Still today: architectures often described with brain terminology (neurons, activations, ...)






Brain as network of information units  
(*neurons*) that can share information






# Single neuron

Sum of incoming connections determines whether neuron will 'fire' (threshold)






# Mathematically



Weights control sensitivity of neuron to information

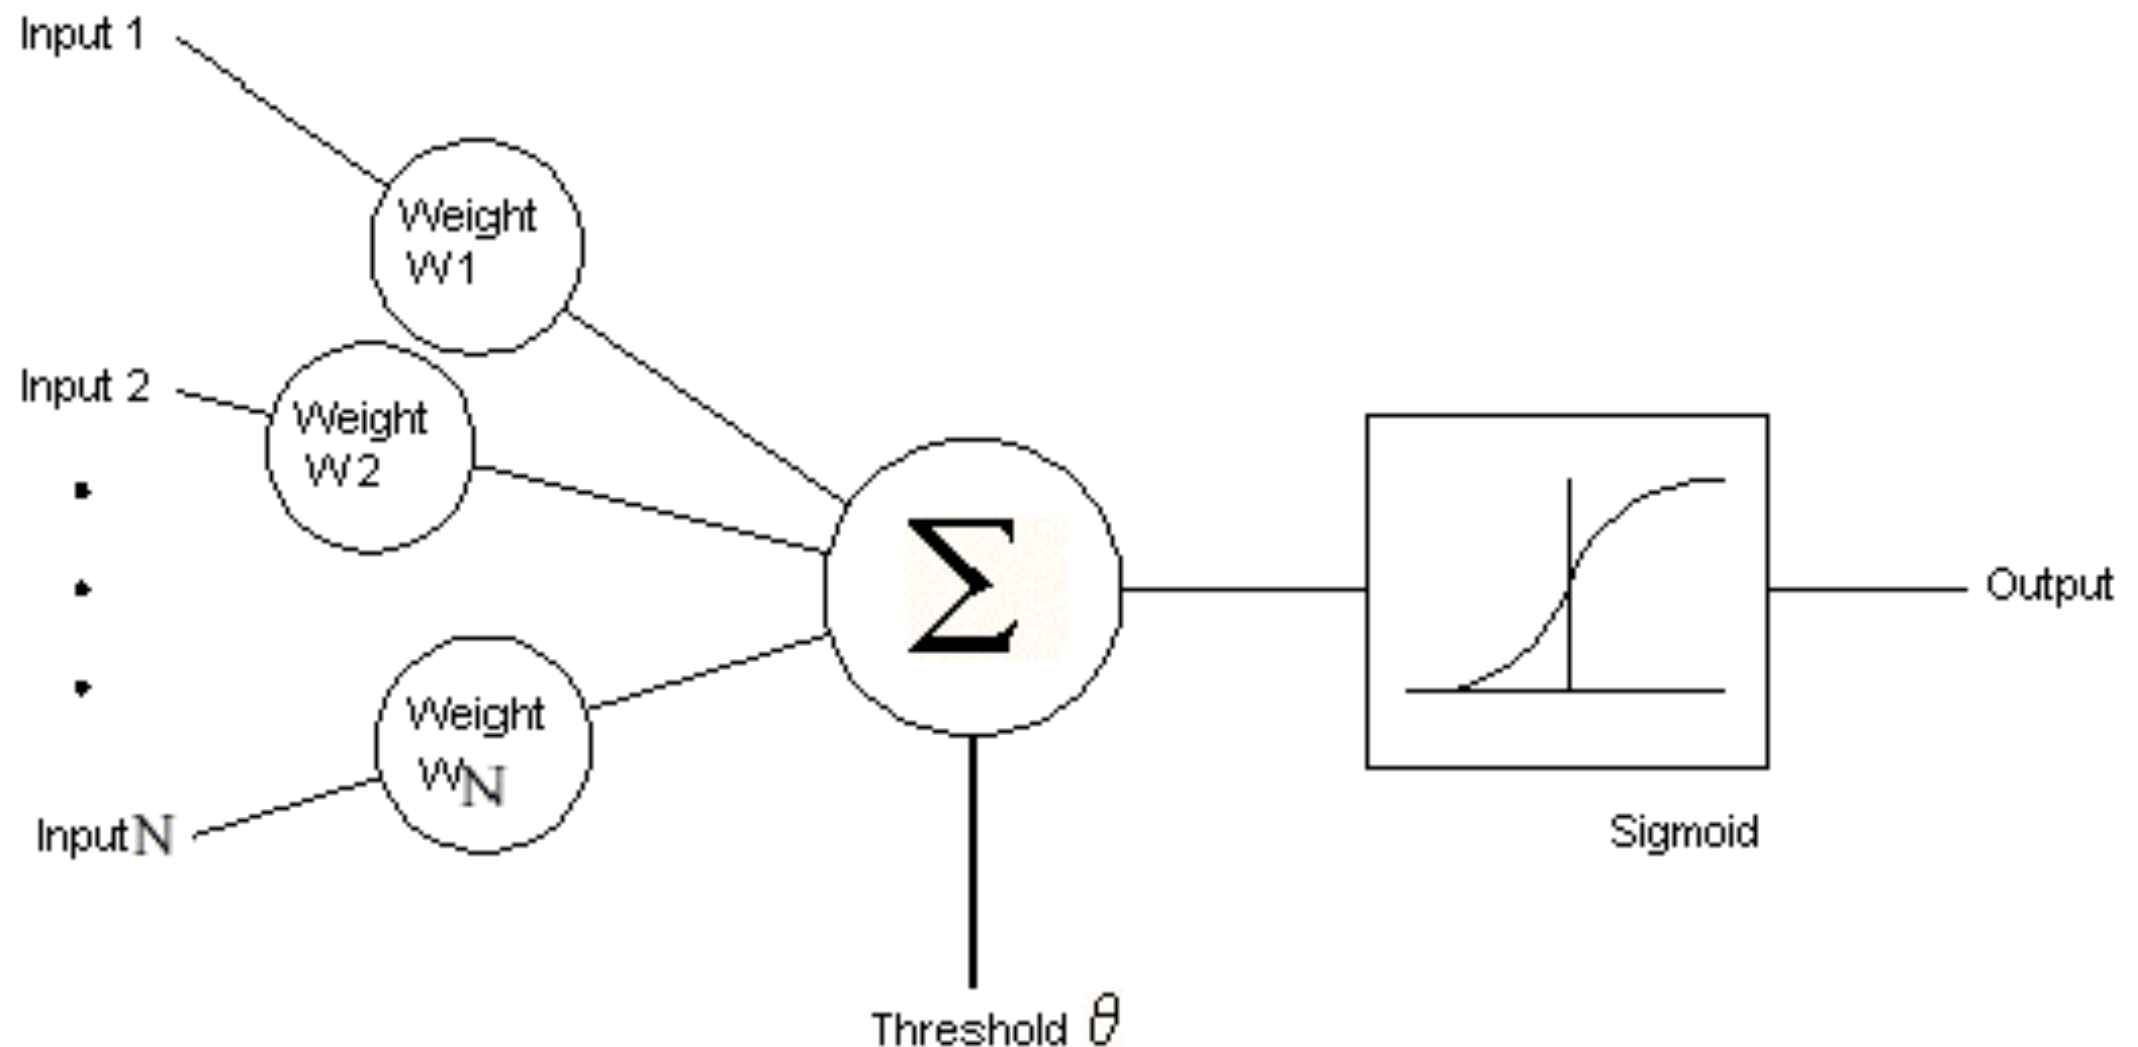
# Activation function

Squash info in range [0, 1]



Sigmoid activation (historically dominant)

# Perceptron



Already useful for regression (*single output*) in ML  
E.g. predict house prices using location, bedrooms, ...



# Notebook: Perceptron

# Training

How do we avoid having to set weights ourselves?

# Tuning an Old Radio



Initially lots of noise, reduce by turning knob

Volume

Frequency




# Tuning an Old Radio



Volume


Frequency





# Intuition?

- Left and right
- Movements get slower as you finetune: learning rate
- You don't know how the radio works internally: only knob and a loss estimate
- Naming conventions:
  - radio = system; knob = parameter
  - sound quality = loss function (which we want to minimize)



# In neural networks?

System or function with many more knobs,  
but exact same principle: one-by-one adjustments







# Training?

- Find the parametrization which minimizes the loss function
- How? Hard, slow and ugly ways:
  - Random search?
  - Try out +/- for each knob and keep best setting
  - ...

# Learning rate



/s real issue in practice (cf. radio):

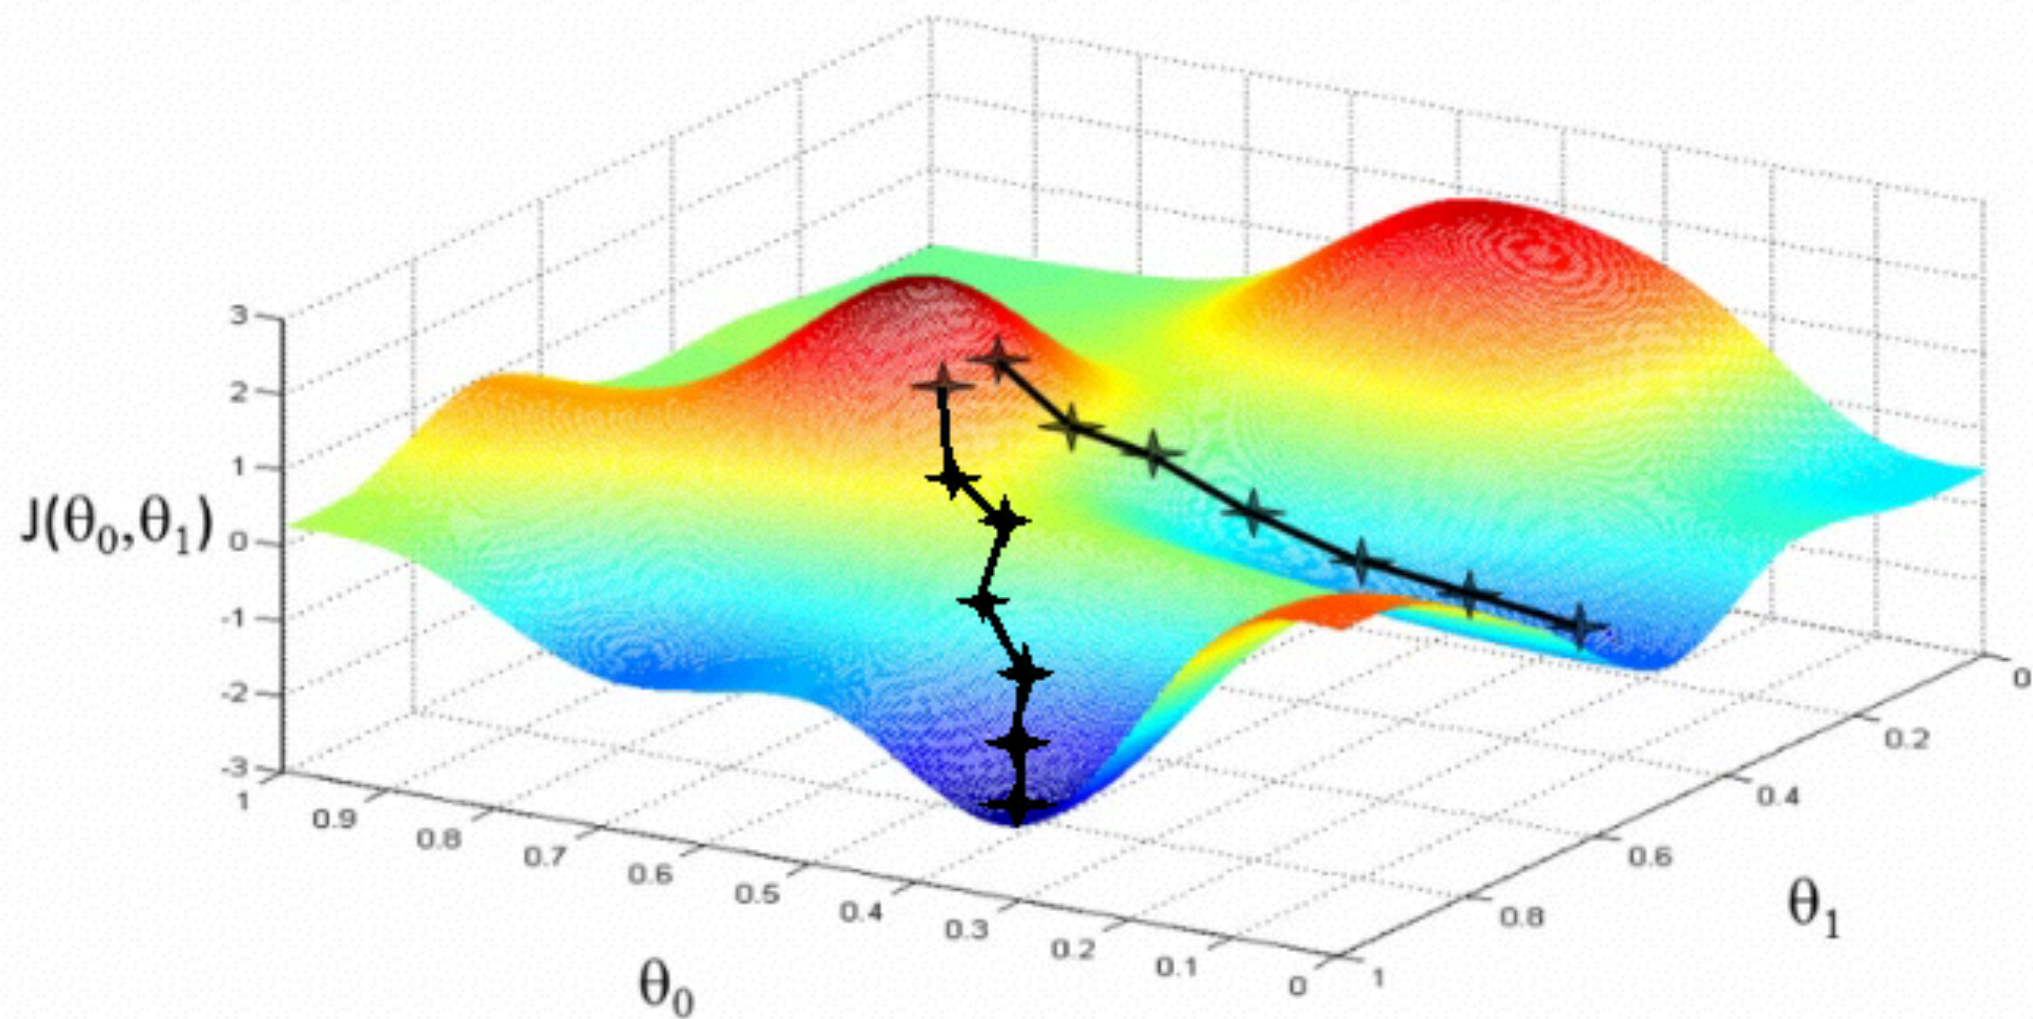
- too **low** a learning rate: convergence to slow
- too **high** a learning rate: you 'miss' the optimum

# Notebook: Naive optimization




# Minimize loss: or 'objective' function

Ski down a hill, preferably ASAP



# What is wrong with our code?

- Our optimisation is (ugly, but esp.) slow:
- calculate results both 'minus' and 'plus'
- Our learning steps are the same throughout



# Solution

- We don't have to calculate plus and minus...
- Because we can calculate the gradient of each parameters!
- Partial derivative
- 'Gradient ascent'

$$\frac{\partial f}{\partial x}$$


'How much does a change in parameter  $x$  affect the total model  $f()$ ?'



# Optimization?

- Calculate partial derivatives for each parameter
- Update each parameter using rule:
  - $\text{param} -= \text{learning\_rate} * \text{gradient}$
- Negative gradient: parameter grows larger

# Solution: Gradient descent



# How?

- Gradient descent to be inflexible:
  - manual derivation
  - hard-coding
- Now: libraries for automatic differentiation (you don't to know the math!)
- Python: Theano, TensorFlow, ...
- You specify  $f()$ : library can return the gradients

theano






# Notebook: gradient descent in Theano

# From simple perceptron....

One class, one weight vector

'weighted sum'  
`np.multiply(X, w).sum()`



# ... to classification?

$n$  classes,  $n$  weight vectors, 1 weight matrix


‘weighted sum’

`np.multiply(X, w).sum()`



‘dot product’

`np.dot(X, W)`





Feature vectors (X)

output

"Dot Product"

The diagram illustrates a dot product calculation. On the left, a feature vector  $X$  is shown as a 2x3 matrix:  $\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$ . The first row (1, 2, 3) is highlighted with a yellow background. In the center, a weight matrix  $W$  is shown as a 3x2 matrix:  $\begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix}$ . The first column (7, 9, 11) is highlighted with a yellow background. A blue 'x' symbol is between the two matrices, and a blue '=' symbol is to the right. A yellow curved arrow labeled "Dot Product" starts from the first row of  $X$  and points to a yellow circle containing the number 58, which is the output. A white arrow points from the text "Feature vectors (X)" to the first matrix, and another white arrow points from the text "output" to the circle with 58. A third white arrow points from the text "Weight matrix (W)" to the second matrix.

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix} = \begin{bmatrix} 58 \end{bmatrix}$$

Weight matrix (W)

# The 'Dense' Layer

A dot product of an input matrix  
with a weight matrix, and  
addition of bias vector


$$output = X \cdot W + bias$$

The **single most fundamental** building block  
in deep learning. All fancy stuff goes back to this!

# Classification



- Use 'softmax' to produce probabilities
- Also [0-1] normalization
- Select class with highest probability
- “Logistic Regression”

$$P(y = j|\mathbf{x}) = \frac{e^{\mathbf{x}^T \mathbf{w}_j}}{\sum_{k=1}^K e^{\mathbf{x}^T \mathbf{w}_k}}$$






# “Deep” Learning?



Stack ‘hidden’ layers between input and output layer

# Computer Vision


## Importance of layers





# Low-level features

Used to be 'handcrafted'!






# Higher-level features



# Analogies human brain



e.g. [Cahieu et al. 2014]

# Notebook: A Multilayer Perceptron in Keras